

Yifan Zhu

📍 Waltham, MA ✉ zhuyifan@brandeis.edu ☎ (202)-758-8938 in linkedin.com/in/yifan-zhu-b67666261
🔖 Google Scholar 🌐 kansusu.github.io

Objective

Computer Science Ph.D. candidate at Brandeis University specializing in **LLM alignment**, **Theory of Mind and epistemic reasoning in dialogue**, and **multimodal common-ground tracking**. Research investigates how humans and AI agents build, track, and repair shared understanding in collaborative interaction — combining formal epistemic modeling, multimodal perception pipelines, and beyond-preference post-training methods. Demonstrated record of leading projects from theoretical framework to deployable system and peer-reviewed publication across NAACL, AAAI, LREC, and EMNLP-track venues. Seeking a **research-oriented Summer 2027 internship** to apply this expertise to production-scale agent systems.

Education

Brandeis University <i>Ph.D. in Computer Science</i>	<i>Expected 2028</i>
◦ Advisor: Prof. James Pustejovsky, Brandeis Lab for Linguistics and Computation (LLC)	
◦ Research areas: Theory of Mind & epistemic reasoning, multimodal perception, LLM alignment, computational semantics	
Georgetown University <i>M.S. in Computational Linguistics</i>	<i>Aug 2019 – May 2021</i>
University College London <i>M.A. in Linguistics</i>	<i>Sept 2015 – Sept 2016</i>
Xiamen University <i>B.A. in Philosophy, Chinese Language and Literature</i>	<i>Sept 2011 – June 2015</i>

Research Experience

Project: Friction-Based Alignment & Theory of Mind	<i>Oct 2025 – Present</i>
◦ Designed and implemented a suite of friction-based LLM alignment methods (FAR, FTR, FPP, GRFR) under the Frictive Policy Optimization (FPO) framework, training agents to issue context-aware interventions conditioned on per-turn epistemic risk rather than static preference signals.	
◦ Introduced a four-identifier Theory-of-Mind schema reconstructing first- and second-order belief states at each referring expression; ablating the second-order channel cuts misunderstanding detection from 65% to 26%.	
◦ Extended FPO from shared-context to perceptually asymmetric dialogue , formulating F^- per-perspective; perspective-bound contradiction detection on MapTask saturates repair-productive friction recall with an open-weights model where omniscient GPT-5 falls short.	
◦ Built the full training pipeline: GPT-5 belief extraction, deterministic friction surrogates, QLoRA fine-tuning of Qwen3.5-27B (DPO, FAR, FPP, FTR), and a calibrated evaluation suite (F1, ECE on warranted / aligned subsets, LLM-judge debiasing).	

Related publications: *Mind the Gap: Theory-of-Mind-Grounded Friction for Epistemic Alignment* (under review, EMNLP 2026, first author); *From Propositional to Perceptual Asymmetry: Extending FPO to Asymmetric Partial Information Dialogue* (under review, SIGDIAL 2026, first author); *Modeling Theory of Mind in Multimodal HCI* (HCII 2024, first author).

Project: Common Ground Tracking & Epistemic Modeling	<i>Mar 2024 – Present</i>
◦ Introduced the Distributed Partial Information Puzzle (DPIP) task and dataset — the first benchmark in which each participant holds a fully disjoint view of a target structure — to evaluate LLM-based and axiomatic DEL-based belief-tracking paradigms.	
◦ Used evidence-based dynamic epistemic logic (EB-DEL) to model beliefs and intentions from multimodal evidence, enabling construction and refinement of common ground in collaborative interactions.	
◦ Linked Generative Lexicon event structures with DEL via Lexical Event Modeling (LEM) , grounding verb meaning in executable pre-/post-state transitions to drive belief updates across multiple modalities.	
◦ Authored multimodal annotation guidelines for common ground and friction labels; led adjudication and annotator training; contributed the block-description transfer module to the real-time TRACE common-ground demo.	

Related publications: Distributed Partial Information Puzzles (LREC 2026, first author); TRACE: Real-Time Common Ground Tracking (NAACL 2025, co-author); Common Ground Tracking in Multimodal Dialogue (LREC-COLING 2024, co-author); Lexical Event Models for Multimodal Dialogues (HCII 2024, co-author); Multimodal Common Ground Annotation Guidelines (2025, lead author).

Project: Multimodal Perception & Situational Awareness

May 2023 – Present

- Developed a **2D perception pipeline** extracting verbal and non-verbal events from audio–video recordings (posture, gesture, object movement, gaze, facial expression).
- Developed a **3D perception pipeline** extracting verbal and non-verbal events from **live-stream data**, enabling real-time analysis of multi-party interactions.
- Designed and implemented an end-to-end **multimodal situational-awareness pipeline** for collaborative problem solving; selected and applied classifiers for CPS codes via **late-fusion** integration of multimodal features.

Related publications: Multimodal Situational Awareness in Collaborative Learning (HCII 2025, first author); Speech Is Not Enough: Interpreting Nonverbal Behavior in Multiparty Dialogue (AAAI 2025, co-author).

Project: Computational Semantics & Cross-Lingual UMR

2022 – 2023

- Conducted sentence- and document-level **Universal Meaning Representation (UMR)** annotation for Chinese, contributing to cross-linguistic semantic analysis.
- Identified Chinese-specific annotation challenges (verb compounds, aspect, anaphora) and proposed compound-annotation solutions adopted in the community guidelines.

Related publications: UMR Annotation of Chinese Verb Compounds (UMR Workshop, 2023, co-author).

Technical Skills

- **Languages:** Python, Bash, LaTeX
- **ML / LLM stack:** PyTorch, HuggingFace Transformers, TRL, PEFT / LoRA / QLoRA, DeepSpeed, vLLM, W&B
- **Alignment & post-training:** RLHF, DPO, GRPO, PPO, reward modeling, KL trust-region control, offline RL
- **Multimodal:** 2D / 3D perception pipelines, speech diarization, gesture / gaze / posture extraction, late-fusion classifiers